

OpenResty TCP 服务代理和动态路由

杭州 OpenResty Meetup

20 December 2017

黄励博(huangnauh)

又拍云

slide



<https://huangnauh.github.io/2017OpenRestyMeetup.html>

<https://github.com/huangnauh/slides>

<http://go-talks.appspot.com/github.com/huangnauh/slides/OpenRestyMeetup.slide>

测试环境

```
git clone git@github.com:huangnauh/slardar.git
git co docker
docker-compose up
```

docker-compose.yml:

```
version: '3.2'

services:
  slardar:
    links:
      - consul
      - mysql3307
      - mysql3306
    image: huangnauh/slardar:test
    volumes:
      - type: bind
        source: ./nginx/conf
        target: /usr/local/slardar/nginx/conf
```

NGINX

NGINX 1.9 开始引入 stream 模块, 实现四层协议的转发和代理, 和 http 类似, 也是采用分阶段处理请求的方式

阶段	简介
Post-accept	接收客户端连接后的第一个阶段
Pre-access	访问的初步检查
Access	实际数据处理之前的客户端访问限制
SSL	SSL 处理
Preread	将数据的初始字节读入预读缓冲区中
Content	实际处理数据
Log	记录请求处理结果的最后阶段


Image credit: [stream_processing](http://nginx.org/en/docs/stream/stream_processing.html) (http://nginx.org/en/docs/stream/stream_processing.html)

SNI 代理

比如, 模块 [ngx_stream_ssl_preread](http://nginx.org/en/docs/stream/nginx_stream_ssl_preread_module.html) (http://nginx.org/en/docs/stream/nginx_stream_ssl_preread_module.html) 在 preread 阶段, 从 ClientHello 消息中提取信息

```
stream {  
    server {  
        listen 443;  
        ssl_preread on;  
        proxy_pass $ssl_preread_server_name:$server_port;  
    }  
}
```

TCP 负载均衡



```
stream {
    upstream mysql {
        #hash $remote_addr consistent;
        server 127.0.0.1:3306;
        server 127.0.0.1:3307;
    }
    server {
        listen 3151;
        proxy_pass mysql;
    }
}
```

TCP 负载均衡

测试:

```
$ echo "show variables where variable_name = 'hostname'" \  
| mysql --skip-column-names -h 127.0.0.1 -P 3151 -uroot -proot
```

```
hostname      e3ac73dd497d <mysql3306 CONTAINER ID>
```

轮询:

```
$ !!;!!
```

```
hostname      396662e2585d <mysql3307 CONTAINER ID>
```

```
hostname      e3ac73dd497d <mysql3306 CONTAINER ID>
```

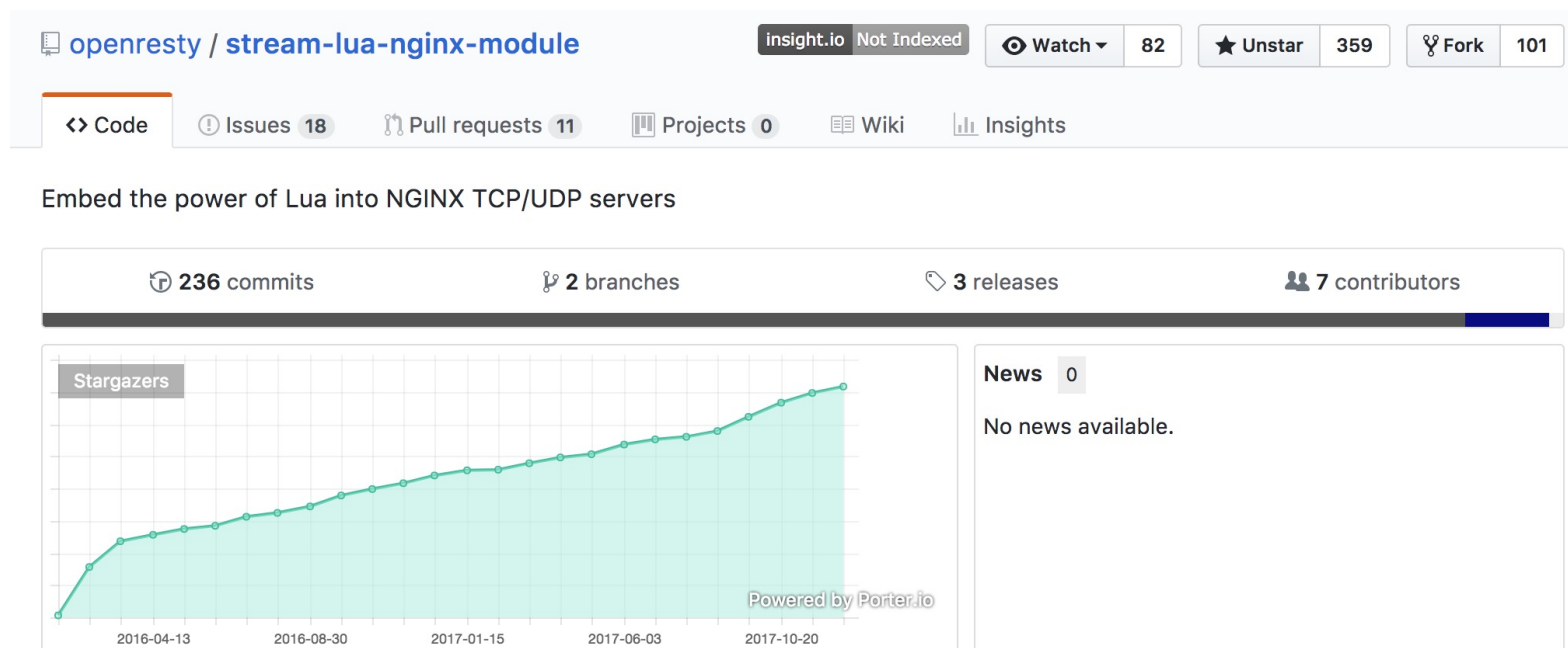
hash:

```
$ !!;!!
```

```
hostname      396662e2585d <mysql3307 CONTAINER ID>
```

```
hostname      396662e2585d <mysql3307 CONTAINER ID>
```

stream-lua-nginx



Description

This is a port of the [ngx_http_lua_module](#) to the NGINX "stream" subsystem so as to support generic stream/TCP clients in the downstream.

Lua APIs and directive names rename the same as the `ngx_http_lua_module`.

[Back to TOC](#)

Hello, Lua!

和 http 类似

```
stream {  
    server {  
        listen 3351;  
  
        content_by_lua_block {  
            ngx.say("Hello, Lua!")  
        }  
    }  
}
```

测试:

```
$ nc 127.0.0.1 3351  
Hello, Lua!
```

TCP 负载均衡

用 Lua 来管理 upstream

```
stream {  
    upstream backend {  
        server 0.0.0.1:4321; # an invalid address  
        balancer_by_lua_file balance.lua;  
    }  
    server {  
        listen 3131;  
        proxy_pass backend;  
    }  
}
```

通过 balancer_by_lua* 和 ngx.balancer 来完成动态选择和重试 upstream

lua-resty-checkups

<https://github.com/upyun/lua-resty-checkups> (<https://github.com/upyun/lua-resty-checkups>)

实现动态 upstream 管理, 之前适用于 http 子系统, 现在也同样适用于 stream 子系统



动态选择 upstream

在 balancer_by_lua* 中, 通过 tcp 的端口来选择相关可用的 upstream

```
skey = ngx.var.server_port

local peer, err = checkups.select_peer(skey)
if not peer then
    ngx.log(ngx.ERR, "select peer failed, ", err)
    return
end

local ok, err = balancer.set_current_peer(peer.host, peer.port)
```

- 其中, peer.host 不支持域名

在测试服务中, set_current_peer 不能直接使用 { host = "mysql3306", port = 3306 }, 需要自己完成解析的操作

checkups 配置

```
_M.["3131"] = {  
  -- 主动健康检查  
  enable = true,  
  typ = "mysql",    -- connect mysql  
  
  -- mysql 信息  
  "user": "runner",  
  "pass": "runner123456",  
  "name": "upyun",  
  
  -- mysql 地址  
  cluster = {  
    {  
      servers = {  
        -- mysql3306 被动健康检查  
        { host = "127.0.0.1", port = 3306, "fail_timeout":10, "max_fails":3 },  
        -- mysql3307 同上  
        { host = "127.0.0.1", port = 3307 },  
        -- invalid  
        { host = "127.0.0.1", port = 3308 },  
      }  
    }  
  }  
}
```

管理 upstream

通过 checkups 的接口, 实现了一个自定义的 tcp 协议, 完成 upstream 的增删改查

Command:

```
<method> <topic> <name>\n
```

```
[ 4-byte size in bytes ][ N-byte data ]
```

DATA Format:

```
[ 4-byte Size ][ 4-byte Type ][ N-byte data ]
```

Type:

OK	0
MESSAGE	1
ERROR	2

管理 upstream

```
$ echo -ne 'PUT upstream 3131\n\x00\x00\x00\x03d{"cluster":[{"servers":[{"host":"127.0.0.1",  
"port":3306}]}]}' | nc 127.0.0.1 1895 | xxd  
00000000: 0000 0006 0000 0000 4f4b                .....OK
```

```
$ echo -ne 'get upstream info\n' | nc 127.0.0.1 1895  
00000000: 0000 02e5 0000 0001 ...
```

info 信息:

```
"3131": {  
  "cluster": [  
    { "servers": [{  
      "host": "127.0.0.1",  
      "port": 3306,  
      "weight": 1,  
      "fail_timeout": 10,  
      "max_fails": 3}]  
    }  
  ]  
}
```

健康检查

```
$ echo -ne 'get upstream status\n' | nc 127.0.0.1 1895 | xxd
00000000: 0000 025a 0000 0001 ...
```

status 信息:

```
{
  -- checkups heartbeat timer is alive.
  "checkup_timer_alive": true,

  -- last heartbeat time
  "last_check_time": "2017-12-20 15:40:58",

  -- status for 3131 cluster
  "cls:3131": [
    [{
      "server": "3131:127.0.0.1:3306",
      "msg": null,
      "status": "ok",
      "lastmodified": "2017-12-20 15:53:21",
      "fail_num": 0
    }]
  ]
}
```


存储 upstream

通过 checkups, 我们可以

- 选择一个工作良好的 upstream
- 对 upstream 进行增删改查

我们还需要: 一个外部数据源来载入 upstream



lua-resty-store

The screenshot shows the GitHub repository page for 'slardar / lua-resty-store'. The repository is owned by 'upyun' and has 35 insights, 240 stars, and 52 forks. It is currently on the 'master' branch. The repository structure is 'slardar / nginx / app / lib / resty / store /'. The latest commit is 'a6dff01' from 'huangnauh' 20 days ago. The repository contains the following files:

File	Commit Message	Time
..		
api.lua	update stream-lua-nginx	25 days ago
config.lua	lua-load in stream-lua	20 days ago
load.lua	lua-load in stream-lua	20 days ago
utils.lua	update stream-lua-nginx	25 days ago

- **api:** consul 和 etcd 的 api kv 接口
- **config:** 从 consul 或 etcd 动态加载配置
- **load:** 从 consul 或 etcd 动态加载 lua 源码

upstream in consul

CONFIG/SLARDAR/UPSTREAMS_STREAM/ +

3131

mysql

config/slardar/upstreams_stream/3131

```
{
  "typ": "mysql",
  "enable": true,
  "user": "runner",
  "pass": "runner123456",
  "name": "upyun",
  "timeout": 2,
  "servers": [{ "host": "127.0.0.1", "port": 3306 }]
}
```

code in consul

CONFIG/SLARDAR/LUA_STREAM/ +

script.preread3151

config/slardar/lua_stream/script.preread3151

```
ngx.exit(1)
```

UPDATE

CANCEL

☐ Validate JSON

DELETE KEY

lua-resty-load

<https://github.com/huangnauh/lua-resty-load> (<https://github.com/huangnauh/lua-resty-load>)

从外部数据源动态加载 lua 源码, 无需 reload nginx



操作 lua 脚本:

与操作 upstream 采用同样的 tcp 协议

```
$ echo -n 'ngx.exit(1)' | wc -c | xargs printf "0x%0.2x"
0x0b
$ echo -ne 'LOAD code script.preread3151\n\0x00\0x00\0x00\0x0bngx.exit(1)' |
nc 127.0.0.1 1895 | xxd
00000000: 0000 0006 0000 0000 4f4b                .....OK
```

测试:

```
$ mysql -h 127.0.0.1 -P 3151 -uroot -proot
ERROR 2013 (HY000): Lost connection to MySQL server
```

获取脚本信息:

```
$ echo -ne 'GET code info\n' | nc 127.0.0.1 1895
{
  "modules": [{
    "time": "2017-12-20 13:54:58",
    "version": "50e9bb007a4a0b3dbd22712f5453a5f1",
    "name": "script.preread3151"}]
}
```

应用举例

流量控制, 以漏桶算法(Leaky Bucket) [resty.limit.req](https://github.com/openresty/lua-resty-limit-traffic/blob/9ac7c27212474ceb20213aea4bbf6c673a009d80/lib/resty/limit/req.md) (https://github.com/openresty/lua-resty-limit-

[traffic/blob/9ac7c27212474ceb20213aea4bbf6c673a009d80/lib/resty/limit/req.md](https://github.com/openresty/lua-resty-limit-traffic/blob/9ac7c27212474ceb20213aea4bbf6c673a009d80/lib/resty/limit/req.md)) 为例:

```
local lim = limit_req.new("stream_limit_req_store", 1, 3)
local key = ngx.var.remote_addr
local delay, err = lim:incoming(key, true)
if not delay then
    return ngx.exit(1)
end
if delay >= 0.001 then
    ngx.sleep(delay)
end
```

现在还不支持 access_by_lua, 在 preread 阶段完成限制功能

加载前:

```
fmt.Printf("connected: %s\n", elapsed)
```

[Run](#)

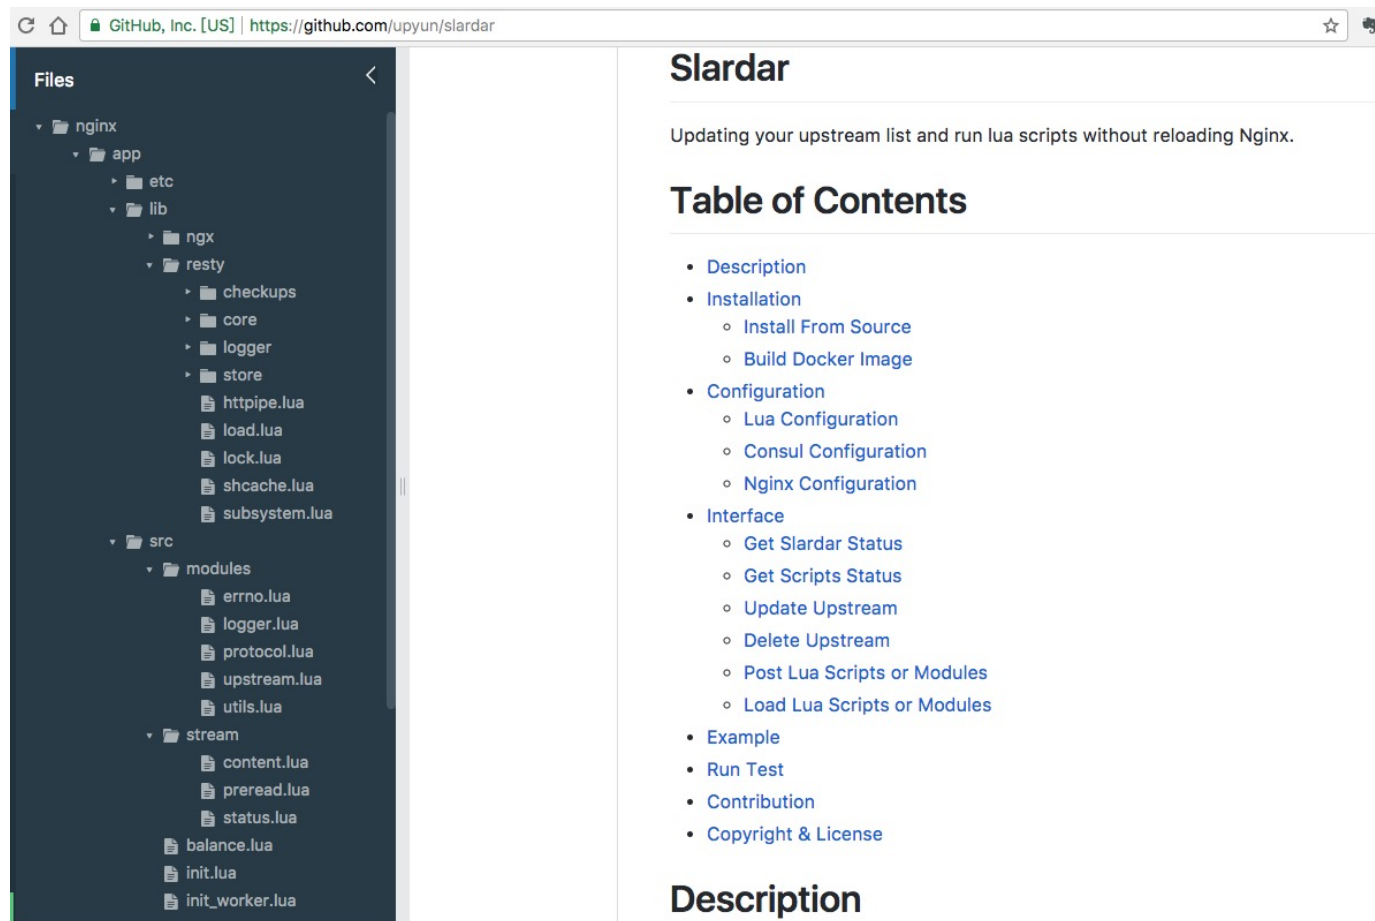
加载后:

```
fmt.Printf("connected: %s\n", elapsed)
```

[Run](#)

Slardar

<https://github.com/upyun/slardar> (<https://github.com/upyun/slardar>)



MySQL Proxy

mysql packet

Payload

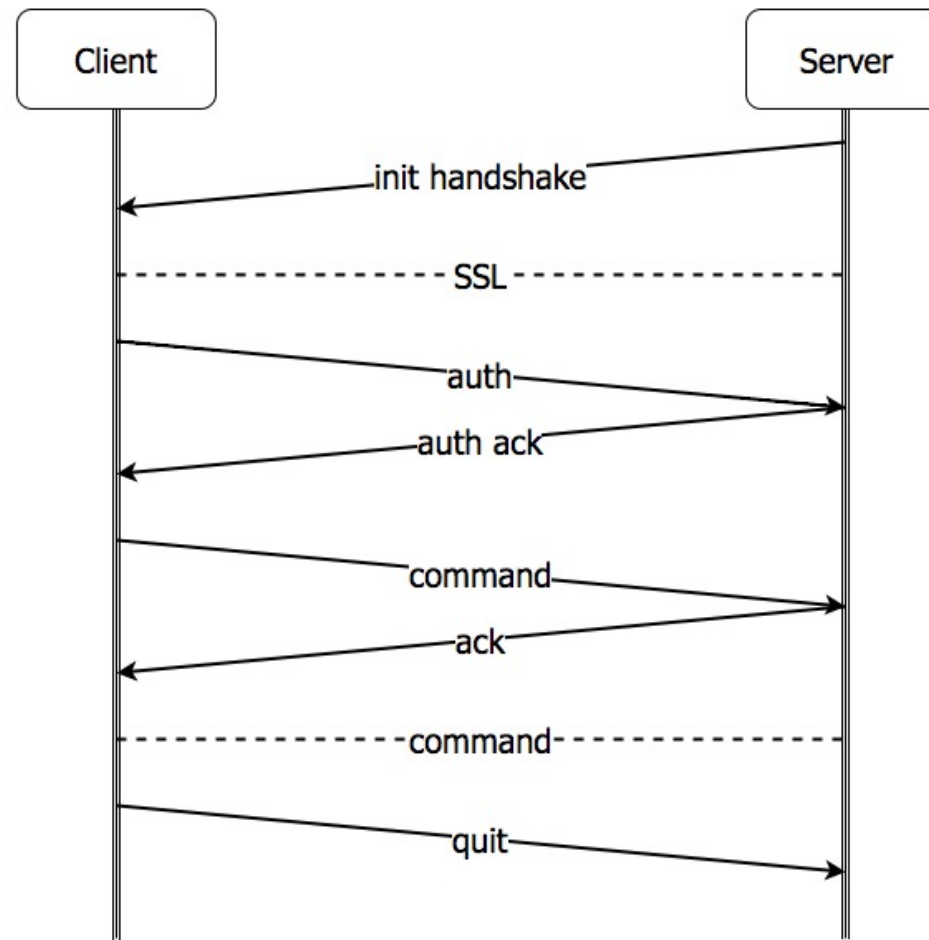
Type	Name	Description
int<3>	payload_length	Length of the payload. The number of bytes in the packet beyond the initial 4 bytes that make up the packet header.
int<1>	sequence_id	Sequence ID
string<var>	payload	payload of the packet

Image credit: [mysql packets](https://dev.mysql.com/doc/dev/mysql-server/8.0.0/page_protocol_basic_packets.html/) (https://dev.mysql.com/doc/dev/mysql-server/8.0.0/page_protocol_basic_packets.html/)

- fixed length integer
- length encoded integer 根据第一个 byte 转换 integer
- null terminated string
- length encoded string 根据开始的 integer 决定 string 长度 (客户端认证数据)

lua-resty-mysql [pr69](https://github.com/openresty/lua-resty-mysql/pull/69) (https://github.com/openresty/lua-resty-mysql/pull/69) 在获取字符串的时候没有把 null terminated string 的 null 去除掉

MySQL 通讯协议



MySQL 握手协议

类型	握手初始化	类型	登录认证
int<1>	协议版本 10	int<4>	客户端权能标志
string	服务器版本	int<4>	最大消息长度
int<4>	connection id	int<1>	字符编码(32 utf8)
string<8>	认证随机字符串(scramble) 第一部分 用于认证	int<23>	填充位
string<1>	填充字节 0x00	-----	-----
int<2>	服务器权能标志 第一部分	string	用户名
int<1>	字符编码(32 utf8)	string	认证数据
int<2>	服务器状态标识, 例如是否在事务或者自动提交模式	string	数据库名称
int<2>	服务器权能标志 第二部分/		
int<1>	scramble 长度 $8 + 13 = 21$		
int<10>	填充位		
string<13>	认证随机字符串(scramble) 第二部分		
string	认证插件名 mysql_native_password		

命令消息

- COM_QUERY 包括 select, update, insert, delete 等

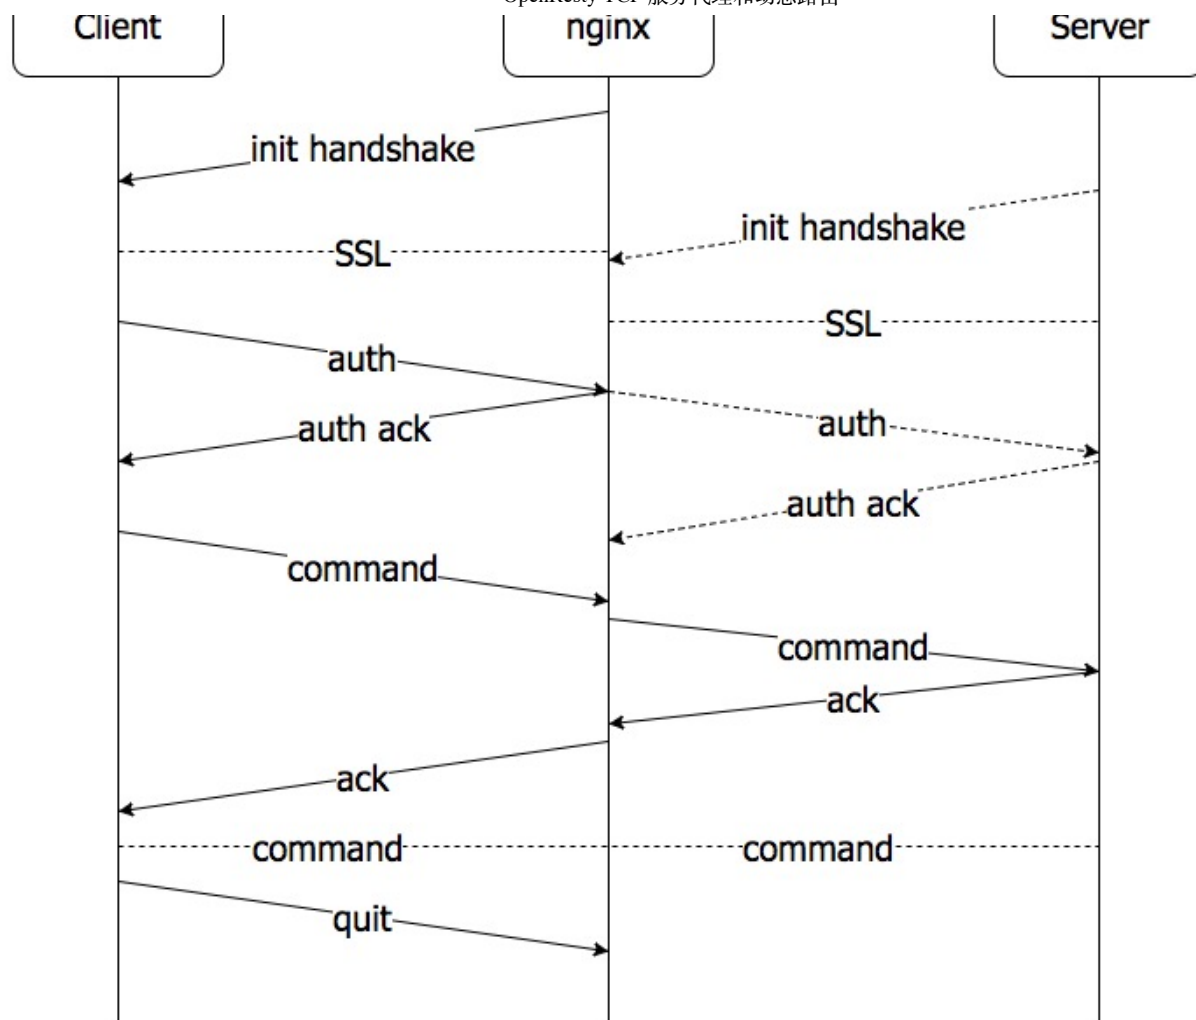
- COM_QUIT 客户端退出

类型	Command	类型	ResultSetHeader
int<1>	命令 COM_QUERY 0x03	int	field_count
string	用户输入	int	extra(可选)

类型	OK Packet	类型	Err Packet	类型	EOF Packet
int<1>	0x00	int<1>	0xff	int<1>	0xfe
int	affected rows	int<2>	错误码	int<2>	告警次数
int	insert id	string<6>	SQL 状态(# 开始)	int<2>	状态标识位
int<2>	服务器状态	string	错误消息		(更多 resultset)
int<2>	0x00				

lua-resty-mysql [pr70](https://github.com/openresty/lua-resty-mysql/pull/70) 不能处理 field_count 大于 250 的情况

MySQL Proxy



lua-resty-mysql-toolset

<https://github.com/huangnauh/lua-resty-mysql-toolset> (<https://github.com/huangnauh/lua-resty-mysql-toolset>)

基于以上介绍的基本协议, 在 lua-resty-mysql 的基础上加入了 server 的协议部分, 包括一个测试用的 proxy

```
stream {  
    server {  
        listen 3141;  
        preread_by_lua_file app/src/stream/preread.lua;  
        content_by_lua_file app/src/stream/content.lua;  
        log_by_lua_file app/src/log.lua;  
    }  
}
```

lua in consul

CONFIG/SLARDAR/LUA_STREAM/ +

script.content3141

config/slardar/lua_stream/script.content3141

```

local proxy = require "resty.mysql.proxy"
local conn, err = proxy:new({user="runner",password="runner123456"})
if err then
    return
end
err = conn:handshake()
if err then
    return
end
conn:process()

```

UPDATE

CANCEL

☐ VALIDATE JSON

DELETE KEY

测试:

```

$echo "show variables where variable_name = 'hostname'" |
pipe> mysql --skip-column-names -h 127.0.0.1 -P 3141 -urunner -prunner123456 upyun

```

```
hostname      huangnauh.local
```

读写分离

1. 通过不同端口来区分读写 upstream, 由应用程序来区分读写

```

stream {
    upstream backend {

```

```
upstream backend {  
    server 0.0.0.1:4321; # an invalid address  
    balancer_by_lua_file balance.lua;  
}  
server {  
    listen 3132;  
    proxy_pass backend;  
}  
server {  
    listen 3133;  
    proxy_pass backend;  
}  
}
```

读写分离

2. 分析 COM_QUERY sql 语句

```
local cmd = string.match(sql, "([^\s,/]+)")  
if not cmd then
```

```
        return nil, "sql error"
    end
    cmd = string.lower(cmd)
    -- 简单 DML 语句区分读写，不考虑带注释的情况
    if cmd == "select" or cmd == "show" then
        ...
    else
        ...
    end
end
```

性能比较

```
sysbench --time=10 --threads=100
```

MySQL Proxy:

```
SQL statistics:
  queries performed:
    read:                48253
    other:               96506
    total:              144759
  queries:              144759 (14396.37 per sec.)
```

MySQL:

```
SQL statistics:
  queries performed:
    read:                65328
    other:              130656
    total:              195984
  queries:              195984 (19525.11 per sec.)
```

Thank you

黄励博(huangnauh)

又拍云

ihuangnauh@gmail.com (mailto:ihuangnauh@gmail.com)

<https://github.com/huangnauh> (https://github.com/huangnauh)

<https://github.com/leandrogilardi/openresty-tcp-proxy> (<https://github.com/leandrogilardi>)

